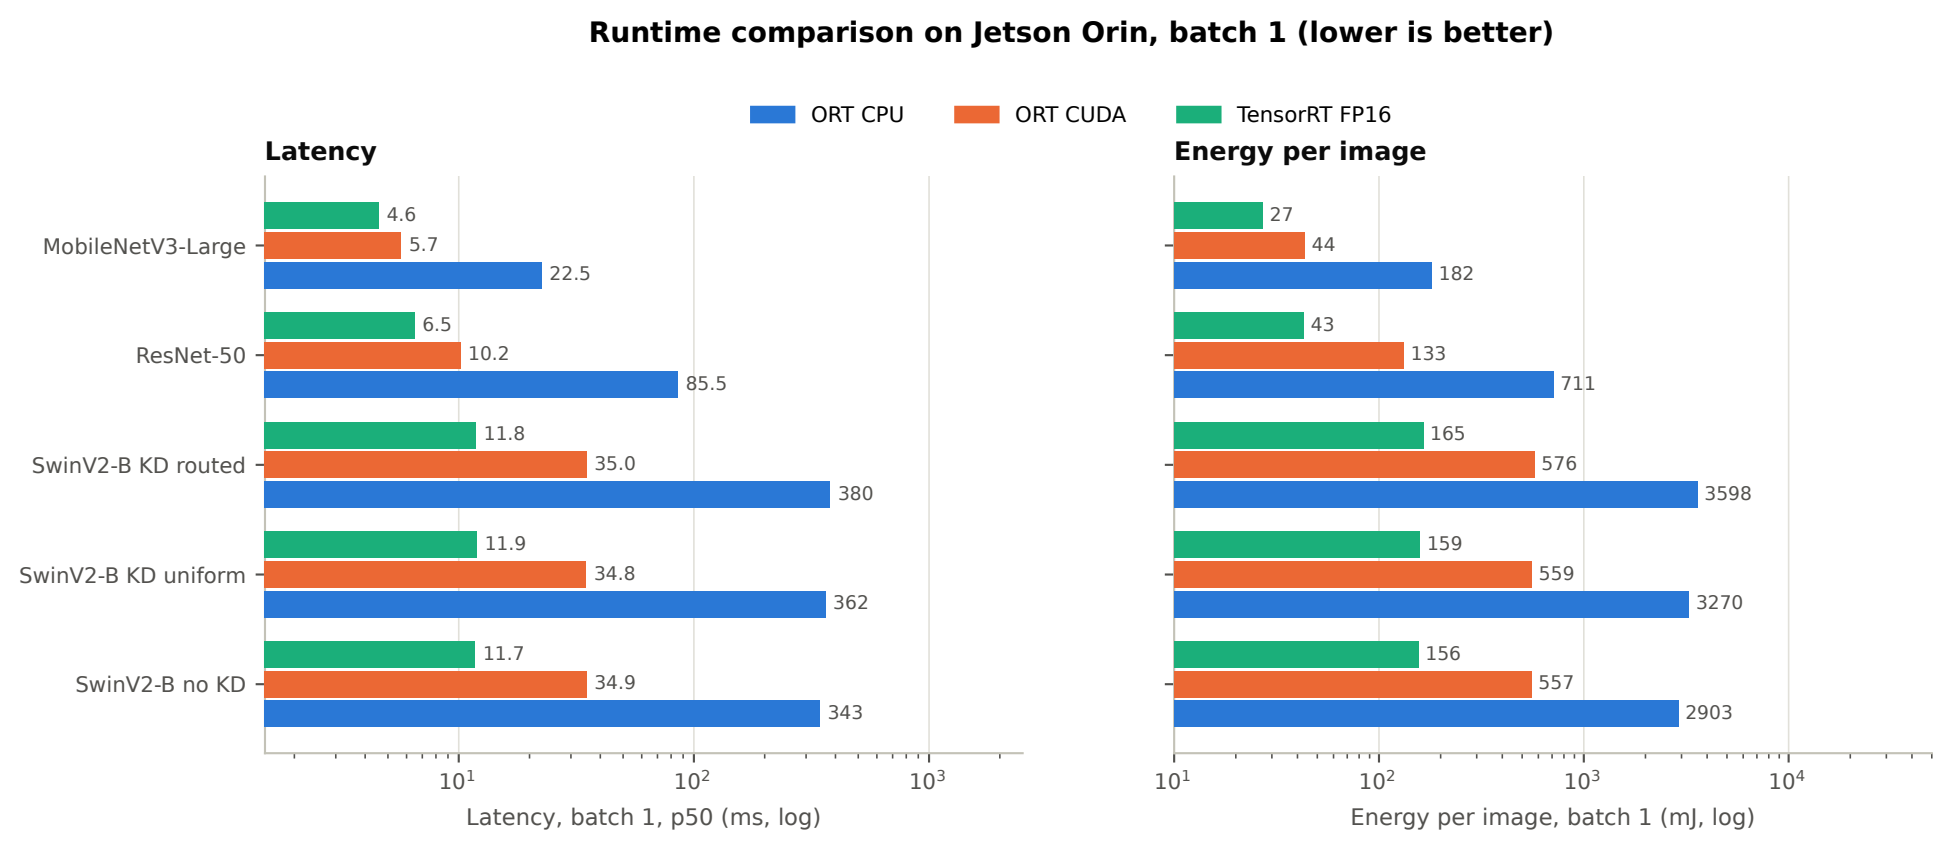
Runtime comparison on Jetson Orin, batch 1 (lower is better)
ORT CPU
ORT CUDA
TensorRT FP16
Latency
Energy per image
MobileNetV3-Large
ResNet-50
SwinV2-B KD routed
SwinV2-B KD uniform
SwinV2-B no KD
4.6
5.7
22.5
6.5
10.2
85.5
11.8
35.0
380
11.9
34.8
362
11.7
34.9
343
27
44
182
43
133
711
165
576
3598
159
559
3270
156
557
2903
10^1
10^2
10^3
10^4
Latency, batch 1, p50 (ms, log)
Energy per image, batch 1 (mJ, log)