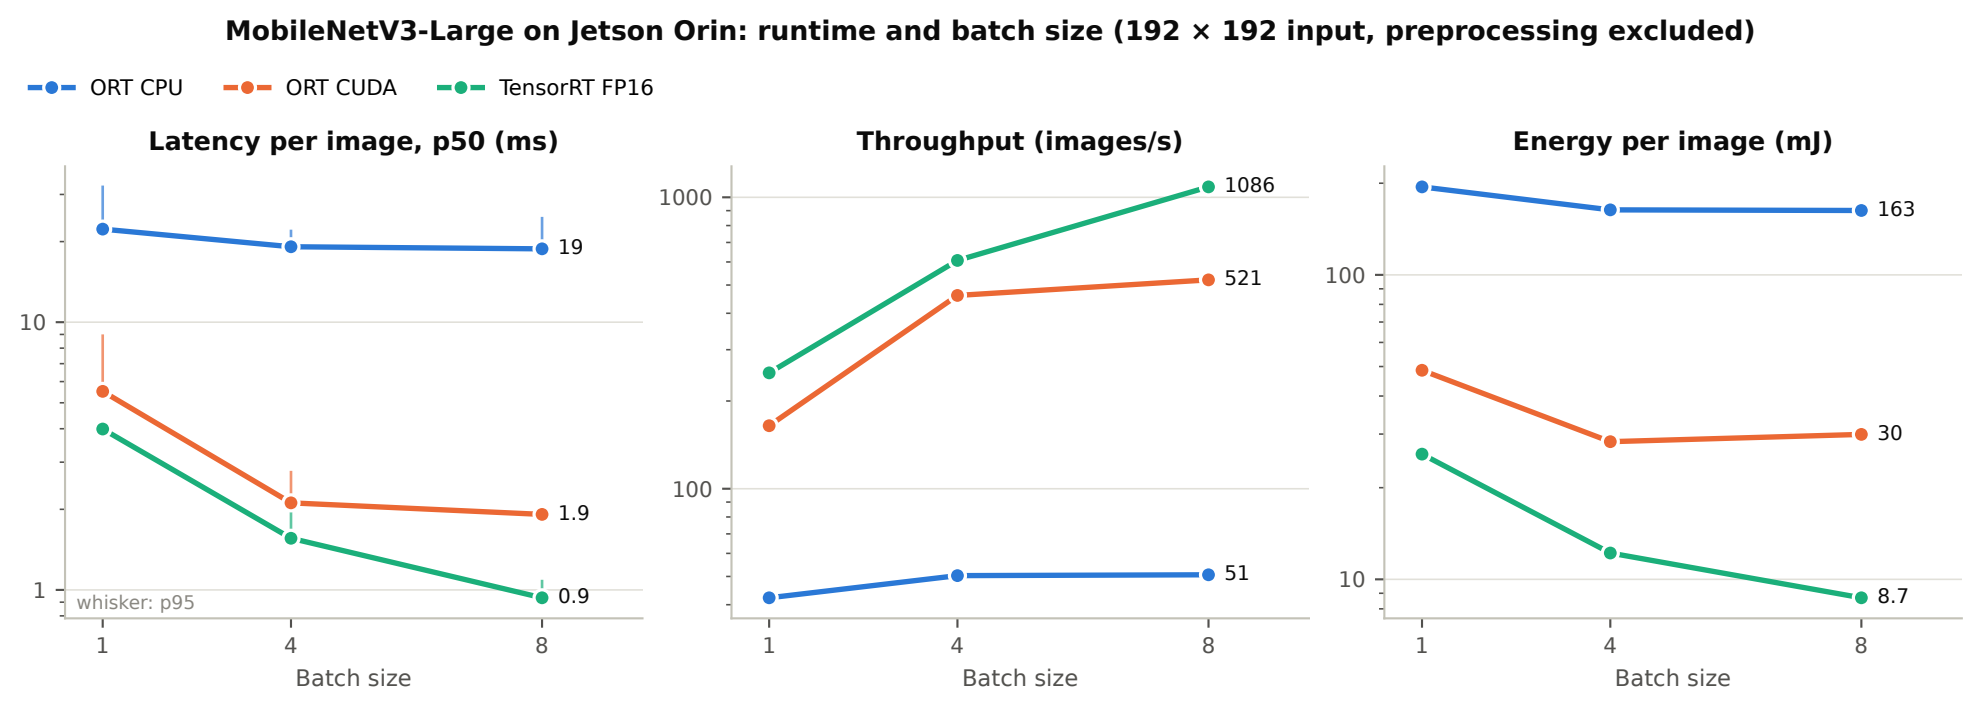
MobileNetV3-Large on Jetson Orin: runtime and batch size (192 × 192 input, preprocessing excluded)
ORT CPU
ORT CUDA
TensorRT FP16
Latency per image, p50 (ms)
10
1
19
1.9
0.9
whisker: p95
1
4
8
Batch size
Throughput (images/s)
1000
100
1086
521
51
1
4
8
Batch size
Energy per image (mJ)
100
10
163
30
8.7
1
4
8
Batch size