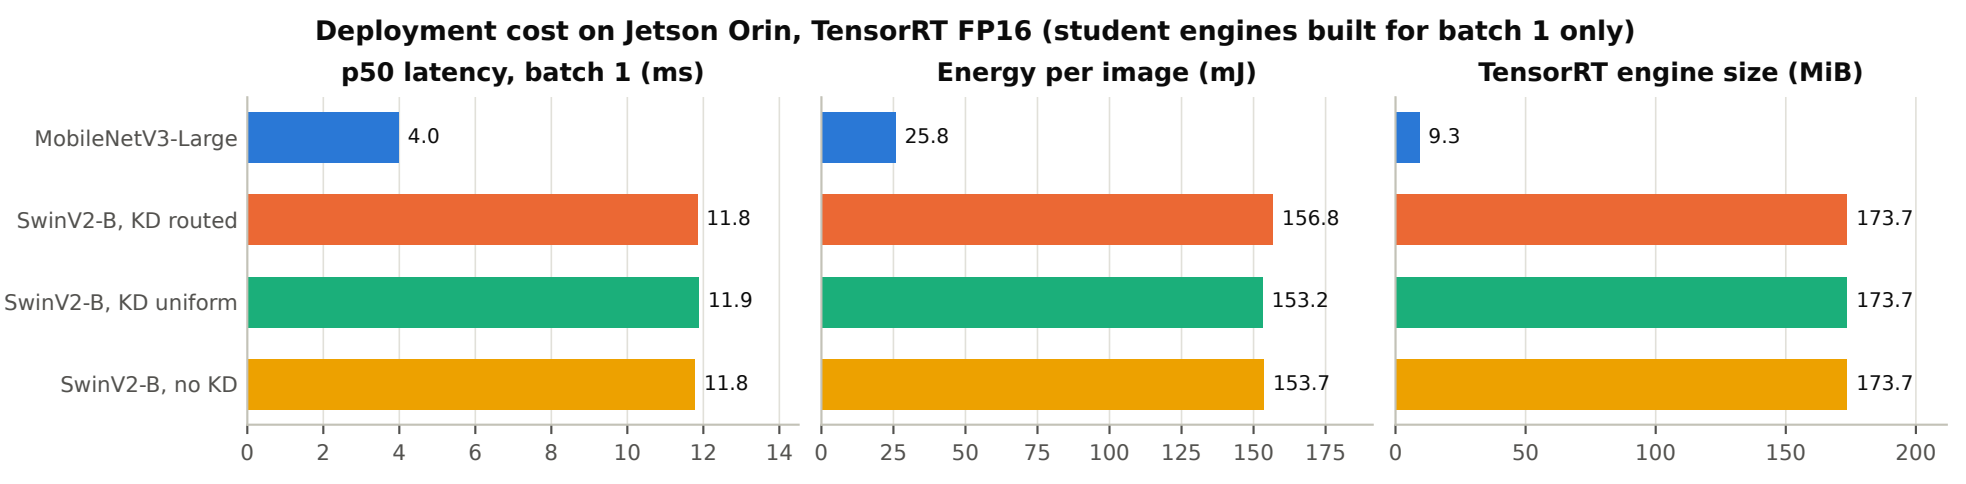
Deployment cost on Jetson Orin, TensorRT FP16 (student engines built for batch 1 only)
p50 latency, batch 1 (ms)
Energy per image (mJ)
TensorRT engine size (MiB)
MobileNetV3-Large
SwinV2-B, KD routed
SwinV2-B, KD uniform
SwinV2-B, no KD
4.0
11.8
11.9
11.8
25.8
156.8
153.2
153.7
9.3
173.7
173.7
173.7
0
2
4
6
8
10
12
14
0
25
50
75
100
125
150
175
0
50
100
150
200