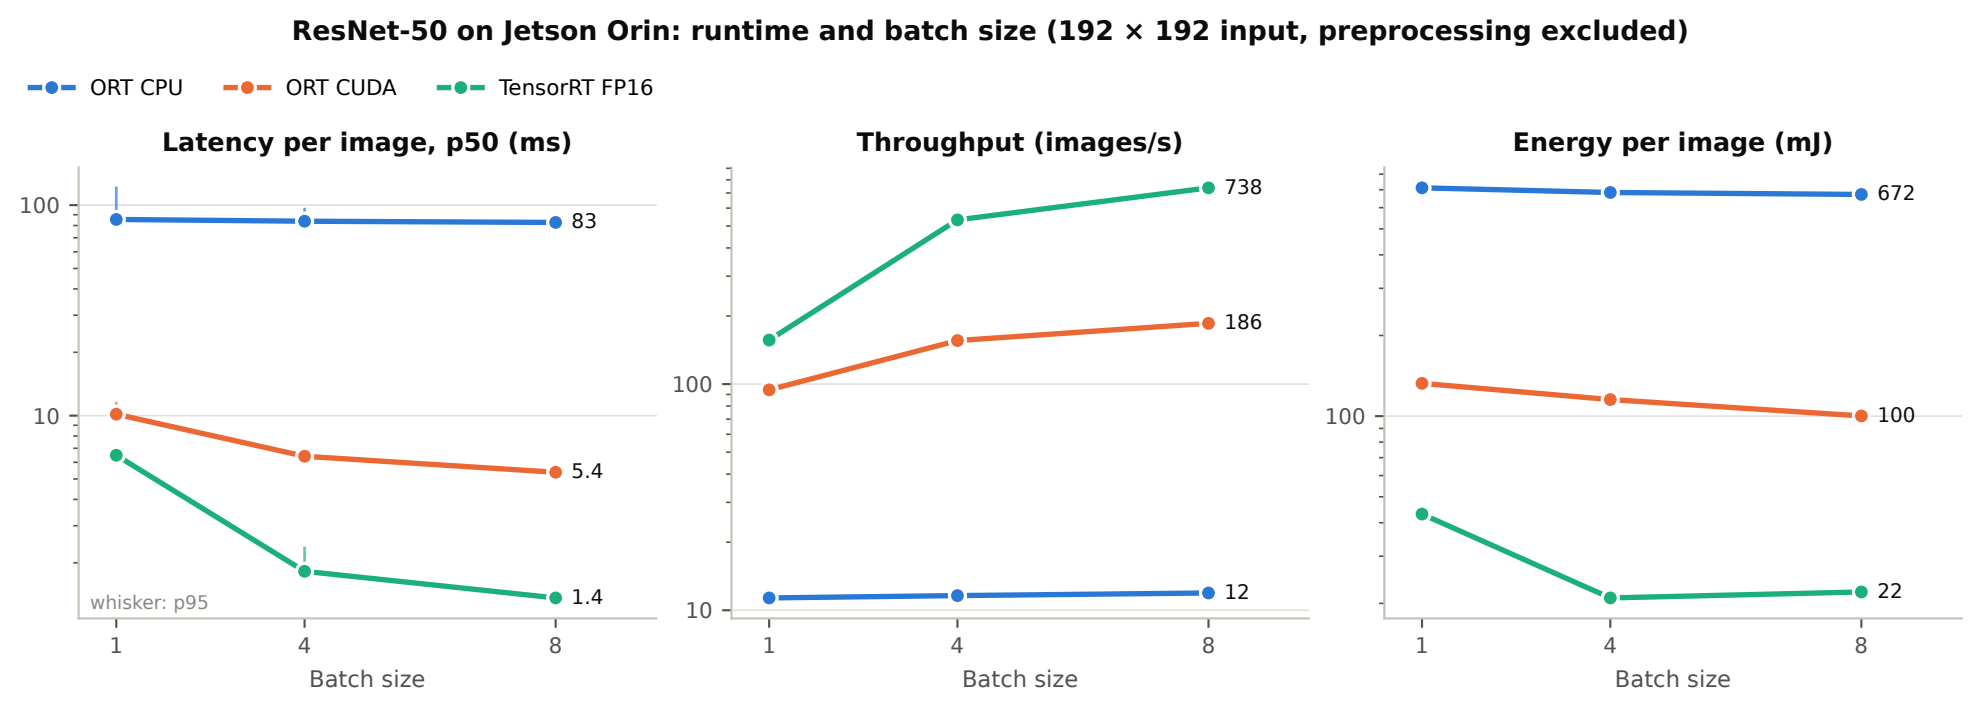
ResNet-50 on Jetson Orin: runtime and batch size (192 × 192 input, preprocessing excluded)
ORT CPU
ORT CUDA
TensorRT FP16
Latency per image, p50 (ms)
100
10
83
5.4
1.4
whisker: p95
1
4
8
Batch size
Throughput (images/s)
100
10
738
186
12
1
4
8
Batch size
Energy per image (mJ)
100
672
100
22
1
4
8
Batch size